

2015 Inter 6fe9d

2015 Internet2 Global Summit

2015 Internet2 Global Summit

- Tuesday
- 8:45 - 10am
- Science Gateways @ Mount Vernon B
 - A science gateway is a community-developed set of tools, applications, and software
 - Mostly instantiated as a web front-end (drupal, django, CMS, etc.) w/ a splash of grouper to back-end compute for sequencing, big data clusters, etc.
 - Client → web interface and REST interface → workbench framework → XSEDE (for parallel) or Triton (for serial)
 - People in US get access up to 50k core hours per year at no cost
 - Galaxy science gateway
 - Focus on reproducibility (many published analyses are not reproducible due to missing software, versions, parameters, etc.)
 - A framework for building gateways
 - Locking to dockerize SW versions to make reproducibility easier
- NSF Proposal Creation @ Mount Vernon A
 - Downloaded presentation
 - Tue 12-1
- Open NTAC Meeting @ Meeting Room 8/9
 - IPv6 WG working on slides and metrics on IPv6 growth on Internet2
 - Performance WG working docs for perfsonar, troubleshooting, measurement/network perf maturity model
 - 3.6 will have more SDN support.
 - Routing and Peering WG some transition of effort re: TR-CPS. Will focus on educating community to make use of TR-CPS, offer tutorials on BGP, etc.
 - SDN no update, they're
 - Security mtg tomorrow at noon. Creation of security WG will be proposed.
 - Topics potentially to be addressed in the near future:

- Wireless
 - Campus
 - Should there be a social responsibility interest group?
 - Around DC, multiple 100gb's now be deployed, particularly for AL2S in DC area
 - OESS doesn't support 100gb lag, doing some shifting of VLAN's around to work around this
 - Not a lot of visibility into traffic patterns from a service perspective, more focused on interface utilization, etc.
 - Path to OpenFlow 1.3 being worked out note that I2's backbone is multi-vendor, so inter-operability is an issue
 - QoS is big driver for OF 1.3
 - Note: fair amount of skepticism around OF 1.3 adoption by hardware vendors
 - That's okay. I2 can use multiple technologies, not just OF
 - Tue 1:15
- Cloud Exploration & Implementation @ Grand Ballroom North
 - Multi university effort to create a cloud cookbook, focused on
 - Federation
 - Case studies
 - Policy and compliance
 - User experience
 - Tue 3pm
- Security Without Borders Meeting Room 12/13/14
 - Downloaded presentation
- PerfSonar Mount Vernon A
 - Been around for 11 years, very part time effort
 - In last 2 years, ESNet, Internet2, Indiana U, and GEANT all started dedicating 1.5 FTE to develop and improve perfSonar
 - 3.5 will be RC this summer and final release before Tech Exchange
 - Should setup perfSonar node as part of RC 100gb
 - perfSonar 3.5
 - Better support for central host mgmt and node auto-config
 - Support for low cost nodes

- Support for other OS choices (RhEL, Debian, VM install)
- Modernize the GUI
- psUI tool from GEANT will be another way to interact w/ perfSonar
- Central mgmt be able to add it to things like puppet, chef, saltstack, cfengine, etc. Sample puppet recipes will be part of 3.5
- New rpm bundles to support this
- Zero config
- Get a DHCP address, register w/ lookup service
- What hosts to run tests to
- Where to send test results
- BWCTL modernized and ported to python, lots of new features added, moved far past its original use case of just being a wrapper for iperf
- Low cost nodes
- Often ARM-based (beagle bone and raspberry pi)
- Can't hit 1gb yet
- Currently targeting around \$200 a node, capable of achieving 1gb BWCTL, allows for Intel
- Best fit for HW
- Gigabyte brix
- Can hit up to 937mbps in testing
- Only draws 30w
- \$225
- ECS Liva
- \$120-\$170
- 32-64gb eMMC, 2gb RAM
- 15w power, could be powered by PoE splitter
- Achieved 940mbps in testing
- Alan W on UH's low cost node deployment
- SWARM dense perfSonar deployment w/ lots of cheap devices (\$50 per node)
- About 50 raspberry pi and beagle bone
- Issues seen
- 22 SD card failures
- 10 file system failures
- 1 corrupted firmware

- Class 10 cards are more reliable
- SanDisk Ultra 8gb class 10
- They liked beagle bone black a lot (sometimes scarce to buy)
- Will do raspberry pi if/when necessary
- Managing them via puppet
- Powstream is really amazing (but boring at first)
- / Perfclub.org/ (perfsonar users group)

Wednesday

- 8:45am Internet2 Advanced Services Today, Mount Vernon A
- Network topology changes
 - Added AL2S nodes
 - Hartford, missoula, indianapolis, cincinatti
 - Backbone upgrades
 - Chi-ash 100gb
 - Ash-mclean 100gb
 - Optical network
 - Hartford add/drop
 - Atlanta metro (56 marietta)
 - MANLAN and WIX
 - Ana-200gb:
 - 100gb Ny to London
 - 100gb Mclean to Paris
 - Layer3
 - 10gb legacy port turndown
 - 3rd 100ge between AL2S and IP in WASH
 - Full mesh between chic, atla, wash, newy, clev
 - TR-CPS
 - Some re-work on west coast and upgrade in Dallas
- Service Updates
 - ION Service Decommission on 4/1
 - NDDI (precursor/test lab to AL2S) network consolidated into lab space

- TR-CPS capacity augments towards peers on west coast, Dallas (imminent!), NY
- General trend - more research traffic on Internet2 in the last year (driven by things like LHC)
- Global Update
 - Response to requests from membership for science and education
 - Atlantic connections
 - ESNet - over 300gb to Europe, across a few circuits
 - Several NREN's have connections as well
 - Nordunet has 100gb between Amsterdam and NY
 - Internet2 has a 100gb as well
 - Completion of a multi-connection ring that doesn't have any single points of aggregation
 - Middle East
 - Open exchange point being built in UAE (Pujarah)
 - Contracted a ckt between Pujarah and Singepore
 - Will connect Pujarah back to London or Amsterdam as well
 - Will likely grow to aggregate many NREN's over the next several years
 - Singapore
 - Colo (10 racks) facility. 6 used, 4 still open
 - Direct links between singapore and HK (several NREN's there)
- Network Security
 - New security role established about a year ago
 - Approach
 - Baseline security risk assessment
 - Develop security program
 - Initiatives
 - Updated security and incident response policies
 - Technical improvements to IDS, strong auth, and vuln mgmt
 - Review access to publically available ops documentation that is sensitive
 - Design and implement a secure OOB mgmt network
 - Build a SOC
 - Guide and improve routing security

- Network Operations
 - Working on ISO 20000 certification
 - 100% availability through investments in infra
 - Strategic NOC plan for 2020
 - Define next gen process automation and predictive analytics
 - Customer driven SLA's
- Network Architecture and Planning
 - Technology
 - Understand tech landscape to build new innovative services, maintain a balance
 - Engagement
 - Engage customers to understand needs, understand vendor capabilities
 - Development
 - Pursue innovative R&D
 - Expand scope of internet's architecture planning to embrace community needs
- Innovation Initiatives
 - Internet2 Service Taxonomy
 - AL1S (fiber and optical transport)
 - NVS virtualized ethernet switching
 - General purpose VLAN Services - AL2S (SDN Controller)
 - Layer 3 services
 - Net +
 - Virtualized Services
 - OESS
 - ONOS
 - GENI Site Mgr
 - Duke's SDN Journey Mount Vernon A
- What is SDN?
 - Control functions decoupled from packet forwarding and switch
 - Controller can view network as a whole
 - Open standards based SW for forwarding decisions (like Openflow)
- Why SDN?

- Firewall and IPS = speed bumps
- Leverage SDN for automated configuration
- Self-service configurable research bypass of core network
- Researcher access to national backbones
- SDN at Duke
 - Implementation of openflow
 - Focused on the edge of the network not datacenter
 - Goal: improve the speed, reliability, and performance of network used by researchers
- Current posture
 - Sdn switches deployed in hub and spoke
 - Switchboard dynamic SDN rule manager
 - SDN controller (modified Ryu REST router)
 - perfSonar nodes deployed across campus
 - Efforts led to redesign of Duke core network
 - Used an MPLS core and can switch to a VRF easily, so routing is every where
- Preparing for SDN
 - Organizational
 - infra and oper readiness,
 - Staffing network, devops, security
 - Culture programmers and systems guys loose on the network
 - Dedicated VRF for SDN control plane
- Infrastructure considerations
 - Dedicated science network?
 - Didn't like this option
 - Converged/unified network?
 - Fiber infra?
 - Needs at the core?
 - Needs at the edge?
- General SDN model at duke
 - Integrated hosts connected to the network

- Hybrid network fabric has multiple options for routing
 - Did not want to build/deploy a separate infra
 - Hosts connected to SDN switches at 10gb
 - Default path for traffic is production network
 - Application applies rules to controller to route certain traffic over alternative path
 - Typically subnet to subnet or host to host
 - VLAN tagging supporting
 - Plan to add functionality (port restrictions, VLAN flipping)
- SDN core network bypass
- Planning for SDN lessons
 - perfSonar is your friend!
 - 10gb cannons aimed at your network!
 - Measure latency and bandwidth on a regular schedule
 - Distribute nodes at multiple locations
 - New version of ps can allow you to split across different interfaces
 - Use ps to prove the network can handle the traffic you expect, BW measurements are useful
 - Use puppet to manage ps
 - Oversubscription
 - Accidental or intentional
 - Span ports, layer 2/3 domains
 - Firewalls
 - What are the real limits per stream / overall?
 - How do they fall over / fail?
 - IPS moved to edge, only IDS and taps in core
- IPS
 - Where is traffic inspected
- It's another network upgrade!
- User-controlled SDN config
 - Once plumbing was in place
 - How did they provide users w/ access?

• Switchboard

- Ruby on rails to simplify SDN config and track changes
- Who is authorized?
- Status of requests
- Update SDN controller based on approved configs
- Rollback/restore SDN controller state
- Audit log of state of configs

• Deployment strategy

- Start with intermittent BW intensive tasks
- Backups
- Bulk data moves
- Data ingestion into protected network segment
- Run building edge switches in hybrid mode
- Enable openflow ports where needed

• Next steps for RYU

- Limit dhcp responses to authoritative dhcp servers
- Vlan tags flipping to support linking VLAN's
- Detect and throttle ARP flooding (and other potential DoS attacks on SDN controller)

• Lessons

- Easy to simulate an SDN network
- Ryu-rest router + mininet + open vswitch
- Open vswitch is very capable
- Openflow port you not care about is disabled
- Bad connections
- How polished is your glass
- Hybrid mode switches in wrong mode

• Endgame

- Campus SDX (software defined exchange)
- Campus core bypass links for science DMZ
- Interconnects layer 2 services (AL2S, BEN, etc)
- Start with a self-service app (switchboard)

- Then enable devops style automation and actions_approvals_audits via switchboard API

- Use cases

- Science DMZ
- AI2s path
- Bypass packet inspection
- Expand BW/capacity
- Data migration (to protect network segment)
- Support Duke/Duke Medicine researchers

- Work in progress

- Mixed support for openflow
- Netflow not necessarily built in

- Staffing_Skills_Culture

- It was easier to teach a programmer to learn networking to run w/ SDN than it was to get a network engineer to learn enough programming/scripting to drive this
- Get network engineers more involved in python scripting, curl etc

#learning/conferences