

### June 2017 Harvard IT Summit

### June 2017 Harvard IT Summit

- Keynote by Nicco Mele, Director Shorenstein Center on Media, Politics, and Public Policy
- 3D printing is awesome!
  - He has been printing lego blocks, t-rex's, even crocs!
  - Then, someone released a blueprint to print a functional handgun
- Thesis technology moves "power" of the institution into the hands of users
- Challenging "at scale"
- People in power often taken by surprise by change
  - World order fragile to technological disruption
- Examples of power moving into hands of the users
  - Personal computer
  - Apple
  - Microsoft Windows
  - Internet
  - Smartphones
- It becomes about:
  - The individual
  - Intimate
  - Intention I search for a white mini-van and google knows that I want to buy a mini-van
- Internet's effect on presidential elections
  - Gary Hart proved he might be able to beat Mondale in the primary in 1984, but took too long to get checks/donations to fund his campaign
  - In 2008, Barack Obama proved he might be able to beat Clinton in the primary and had a website to immediately get donations and fund his campaign
  - Tea Party

- Donald Trump and tweeting unconstrained by norms and institutions
- Major milestones in US history mediated by news/journalism
  - JFK shot
  - Watergate
  - Berlin Wall
  - 911
  - Common experience created by journalism
- Compare these experiences to Bin Laden getting killed
  - Very first announcement a tweet from Keith Urbhan, goes viral, changes White House strategy on announcement
  - Happened w/out news media
- Journalism is collapsing and fake news is winning
- What do we do about it?
- Re-think our relationship
  - Authority and expertise not taken at face value
  - Self-funded solutions to municipal problems normally covered by taxes
- Speed and intensity of response and mutual support after Boston Marathon bombing
  - Donations
  - Google spreadsheet for temporary housing
- Internet also tried to find the 2 bombers
  - When people get angry, they are going to use the Internet as an outlet for their anger
- Demand technical literacy from our leaders and each other
- Fight "nerd disease" technical elitism
  - "genius bar"
  - Pebkac
- Demand accountability One
  - 10:50am Concurrent Session 1 OpenNebula at FAS RC Sever 102
- Presented by:

- Justin Riley, John noss, Wess Dillingham
- Dr Ignacio Llorente Project Director at OpenNebula and visiting scholar at Harvard
- FASRC VM's are provisioned similar to HUIT ticket request, request reviewed if viable, deploy it, make sure it's up and running
- Why Private Cloud?
  - Workload requirements
  - Performance, latency
  - Security, confidentiality, privacy (isolated infrastructure)
  - Leverage existing infrastructure and know-how
  - Storage, networking, config mgmt
  - Cost analysis
  - Variability of demand (x10 cheaper storage)
- The old system based on KVM and GlusterFS
- New System
  - OpenNebula + KVM + Ceph Cluster
  - Multi-datacenter hypervisor and storage
  - Automated VM scheduling
  - Image based OS installs
- Why OpenNebula?
  - Light and simple
  - Robust production ready, highly scalable
  - Flexible
  - Powerful innovative functionality for private/hybrid cloud and DC virtualization
  - 1 week for OpenNebula versus months to really get OpenStack working
- Why Ceph?
  - Killer storage solution for private cloud
  - Software defined storage
  - Acquired and maintained by RedHat
  - File, block, and object store
  - Can backend into S3 and Swift
  - Can scale to exabyte

- RC primarily uses it for block storage for VM hosting storage
- Architecture Review of the full stack
  - They're using OpenNebula to also start provisioning VM's in AWS and Azure
  - Single dashboard for mgmt
- Integration and Mgmt
  - Puppet
  - Using Puppet to do operations and configuration mgmt for OpenNebula
  - Can spin up multiple clusters w/ their puppet scripts
  - Backups
    - Production Ceph Cluster daily disk snapshots
    - Backup Ceph Cluster
    - Deltas between today and yesterday transferred from prod ceph cluster to backup ceph cluster
    - More CPU intensive work done on the backup cluster
    - QCOW2 Deep Backups
    - Export RBD devices from backup cluster to qcow2 file format and keep in separate filesystem in separate datacenter
  - Monitoring
    - Ceph dashboard
    - Web dashboard Grafana
    - Nagios alerting
    - VM Leaderboard which VM's are hammering disk?
  - OpenNebula hooks/plugins
    - IP address check to ensure
    - OneDNS
      - Provide dynamic DNS resolution for VM's on OpenNebula
      - Dynamic DNS generation for all VM's based on sanitized name
      - Automatic forward and reverse records per VM
- Operations
  - They use an internally developed tool called Cangallo to generate clean VM images similar to dockerfile for containers
  - Simple YAML file format

- They use a single git repo for all our base and derived images
- NESE will be based on Ceph
  - 11:45am Lunch
  - 1:00pm Concurrent Session 2 JupyterHub on AWS - Demba and Farras Sadek
- Bridging the gaps between EE and CS education
  - EE data set and data collection
  - CS Data processing and IoT
- Goal integration of theory and computation and design of seamless coding interface
- Jupyter Notebooks on AWS
  - Minimize lead time to get a useful programming environment up and running a "sandbox"
- JupyterNotebook
  - A web application for coding, documentation, simulation
- JupyterHub
  - A way to give a Notebook envier to each person or group of people (PaaS)
  - Users can log in w/out any setup and start usign it
- Hub functions
  - Authentication, authorization,

Manages AWS resources

- 2:20pm Concurrent Session 3 Video Delivery and Retention at Harvard
- Video at HBS Kaltura based
  - Use cases both academic (course videos, etc.) and administrative (marketing, external relations, initiatives)
  - Academic Solution Design
  - KMC Media Console
  - Classroom Delivery KMS (MediaSpace)
  - Playlist
  - LMS
  - CCTL KMS

- Sharepoint
- Administrative Solution Design
- KMC
- KMS
- Sharepoint
- Websites
- Links
- Upfront metadata design and video migration
- Canvas and Kaltura were rolled out at the same time to ease migration disruption
- Extensive time spent w/ faculty on tool requirements gathering and training
- Video Retention and Metadata Recommended Guidelines
  - Video storage workgroup
  - Tasked w/ developing a centrally managed service to manage, distribute, and retain Harvard's video content
  - Defining standards and guidelines metadata and retention in particular
  - Retention and Metadata Sub-groups
  - Recommendations to serve as basis for wider retention and metadata efforts moving forward
- Questions:
  - For donation of lecture capture, Intellectual property owner can choose to donate, which bypasses the default retention policy. What about department head? (lectures that only belatedly recognized for historical significance)
  - For metadata and digital asset management, how does versioning play into asset mgmt?
- Video Metadata Recommendations minimum required fields
  - Title
  - Description
  - Ingest data
  - Creation date
  - GRS Category
  - Retention Requirements
  - Organizational Owner
  - Usage Class

- Filename/path
- Security Level

#learning/conferences